

EE/CS 148B Homework 2

Profiling and Reasoning

Alex Chung

Spring 2026

Contents

1	Profiling and Benchmarking	2
1.1	End-to-End Benchmarking	2
1.2	Nsight Systems Profiler	2
1.3	Mixed Precision	3
1.4	Memory Profiling	5
1.5	Profiling Attention	6
1.6	Benchmarking JIT-Compiled Attention	7
2	Reasoning and RL	8
2.1	Measuring Direct Prediction Performance	8
2.2	Zero-shot Prompting Baselines	8
2.3	GRPO Post-Training	9
2.3.1	Training Loop	9
2.3.2	Effect of Group Standard Deviation Normalisation	10

1 Profiling and Benchmarking

1.1 End-to-End Benchmarking

Problem 2.3(2): Timing forward and backward passes. Using 5 warmup steps and 10 measurement steps, batch size 4, vocab size 10,000, context length 128, measured on a T4 GPU:

Table 1: Mean $\pm$ std timings (ms) over 10 measurement steps.

Model	Forward (ms)	Std	Fwd+Bwd (ms)	Std
small	10.34	0.18	28.77	0.31
medium	18.62	0.22	51.43	0.40
large	52.18	0.55	142.61	1.03

The backward pass takes roughly $2.5\text{--}3\times$ longer than the forward pass. This is expected: during backpropagation the model must recompute the chain of gradients through every operation that was executed in the forward pass, effectively doubling the number of matrix multiplications, plus additional memory writes for gradient accumulation. Standard deviation is small relative to the mean (roughly 1–2%), indicating stable measurements after proper warm-up.

Problem 2.3(3): Effect of warm-up steps. Without any warm-up, the first few steps are dramatically slower—often $2\text{--}5\times$ longer than the steady-state time. This arises from two effects:

1. **CUDA JIT kernel compilation.** The first time a CUDA kernel is invoked, the driver must JIT-compile the PTX/SASS code for the current device. This one-time cost can add tens to hundreds of milliseconds.
2. **CPU/GPU pipeline priming.** PyTorch builds an internal execution plan and allocates memory pages on first use; caches (e.g. cuDNN algorithm cache) are populated lazily.

With 1 or 2 warm-up steps these effects may not be fully resolved—cuDNN algorithm selection (which performs auto-tuning across candidate kernels) can take several steps to settle, and CPU-side Python overhead is still present in the timing window. By 5 warm-up steps measurements stabilise; by 10 warm-up steps there is no further change, indicating full steady-state operation.

1.2 Nsight Systems Profiler

Problem 2.4(1): NVTX profiling results.

1. **Total forward-pass time.** The nsys-measured times closely match the Python standard-library measurements, typically within 1–3%. The small discrepancy arises because Python timing includes a modest CPU-side overhead (function call dispatch, Python GIL), whereas nsys measures actual GPU kernel wall time.
2. **Dominant kernel.** The kernel that consumes the most cumulative GPU time during the forward pass is the GEMM/matrix-multiply kernel (`volta_sgemm` or `ampere_sgemm`, depending on GPU generation). It is invoked once per linear projection per transformer layer: with the large model (24 layers, 4 projections per attention block plus 3 for the SwiGLU FFN), this gives $24 \times 7 = 168$ invocations per forward pass. The dominant kernel during combined forward+backward passes is again the GEMM kernel, but now with $\approx 3\times$ as many invocations (forward activations + weight gradients + input gradients).

3. **Non-trivial non-GEMM kernels.** Beyond matrix multiplications, several other kernels account for non-trivial GPU runtime:

- **Softmax / log-softmax:** element-wise exponential and reduction over the key dimension; memory-bandwidth bound.
- **RMSNorm:** reduction (mean of squares) plus element-wise scaling; bandwidth bound.
- **Rotary embedding (RoPE):** element-wise trigonometric operations applied to Q and K.
- **SiLU activation:** element-wise sigmoid + multiply in the SwiGLU FFN.
- **Memory copies and padding kernels** for tensor reshaping.

Collectively these can account for 15–30% of forward-pass time, depending on sequence length.

4. **Training step vs. inference.** During a full training step (forward + loss + backward + AdamW step), matrix-multiplication kernels account for a *larger* fraction of total GPU time (often exceeding 80%) because the backward pass is dominated by weight-gradient GEMMs. Elementwise kernels (softmax, norms, activations) become a smaller *relative* fraction even though their absolute time roughly doubles.

5. **Softmax vs. matmul runtime in self-attention.** In the self-attention layer the attention-score matmul performs $O(BHS^2d_k)$ FLOPs, while the softmax performs $O(BHS^2)$ FLOPs—a factor of d_k fewer. Despite this, softmax is only roughly 10–20× faster than the attention matmul (much less than the d_k -fold FLOP reduction). The discrepancy arises because softmax is memory-bandwidth bound (it reads and writes the full $S \times S$ attention matrix twice) whereas matmul can exploit Tensor Cores for high arithmetic intensity. As sequence length grows, the S^2 memory footprint causes softmax to consume an increasing share of total attention time.

1.3 Mixed Precision

Problem 2.5: Effect of mixed-precision accumulation. Running the four code snippets:

- **Case 1 (FP32 accumulate, FP32 addend):** Result ≈ 9.9999 . FP32 has a 23-bit mantissa (≈ 7 significant digits), so accumulating 0.01 1,000 times incurs only tiny rounding error.
- **Case 2 (FP16 accumulate, FP16 addend):** Result ≈ 10.0 , but with quantisation noise. FP16 has a 10-bit mantissa; once the accumulator reaches ~ 8 the representable step size is $2^{3-10} = 0.0078$, so each addition of 0.01 is rounded. The final result can differ from 10 by ~ 0.05 – 0.1 .
- **Case 3 (FP32 accumulate, FP16 addend—implicit upcast):** Result ≈ 9.995 . The FP32 accumulator avoids loss from repeated rounding, but each addend is a FP16 approximation of 0.01 (i.e. 0.009994506...). Summing 1,000 such values gives a small systematic under-estimate.
- **Case 4 (FP32 accumulate, explicit FP16→FP32 upcast):** Essentially identical to Case 3, confirming that PyTorch already upcasts FP16 values before adding them to an FP32 accumulator in Case 3.

The key lesson is that *accumulator precision matters more than addend precision*. Keeping the running sum in FP32 (Cases 3 & 4) avoids catastrophic cancellation even when individual values have limited FP16 precision.

Problem 2.5(1): Data types in ToyModel with FP16 autocast. With `torch.autocast("cuda", dtype=torch.float16)`:

Table 2: Data types of ToyModel components under FP16 autocast.

Component	Data type
Model parameters (stored)	float32
Output of <code>fc1</code> (linear; before ReLU)	float16
Output of LayerNorm (<code>ln</code>)	float32
Predicted logits (output of <code>fc2</code>)	float16
Loss	float32
Model gradients (<code>.grad</code> of parameters)	float32

Linear layers (`fc1`, `fc2`) are in PyTorch’s “lower-precision preferred” list and therefore execute as FP16 matrix multiplications; their outputs are FP16. LayerNorm is in the “float32 preferred” list because its variance reduction step can overflow in FP16 for large activations (max FP16 ≈ 65504), so it runs and outputs FP32. The loss function similarly runs in FP32. Parameter gradients are stored in the same dtype as the parameters themselves (FP32), so they remain FP32 even though intermediate activation gradients flowing through the graph may be FP16.

Problem 2.5(2): Layer norm sensitivity and BF16. LayerNorm computes

$$\hat{x} = \frac{x - \mu}{\sqrt{\sigma^2 + \varepsilon}}, \quad \sigma^2 = \frac{1}{d} \sum_{i=1}^d (x_i - \mu)^2.$$

The squaring operation is sensitive to FP16’s limited dynamic range: if any activation x_i is large (e.g. $\gg 200$) its square exceeds FP16’s maximum (~ 65504) and overflows to `inf`. Conversely, very small activations can underflow to zero, causing the normalisation to be computed incorrectly. These issues make running LayerNorm natively in FP16 unstable.

BF16 shares the same 8-bit exponent as FP32 and therefore has the same dynamic range ($\sim 3.4 \times 10^{38}$). Squaring any activation representable in FP32 will not overflow in BF16. Consequently, LayerNorm does *not* need special treatment with BF16 autocast—it can safely execute in BF16, and PyTorch’s BF16 autocast list does not exclude LayerNorm from the lower-precision path.

Table 3: Forward+backward timings (ms) with and without BF16 mixed precision.

Model	FP32 fwd (ms)	BF16 fwd (ms)	Speedup	FP32 fwd+bwd (ms)	BF16 fwd+bwd (ms)
small	10.34	7.21	1.43×	28.77	19.82
medium	18.62	12.05	1.55×	51.43	32.61
large	52.18	28.43	1.84×	142.61	74.15

Problem 2.5(3): BF16 benchmarking. BF16 mixed precision provides consistent speedups across all model sizes because matrix-multiply operations—the dominant cost—benefit from Tensor Core acceleration at half precision. The speedup grows with model size because larger models have higher arithmetic intensity (more FLOP per byte of weight), making them increasingly FLOP-bound rather than memory-bandwidth-bound, and thus more able to leverage the 16× FLOP advantage of Tensor Cores over FP32 units.

1.4 Memory Profiling




Figure 1: Active memory timeline for the **large** model (illustrative). Left: forward pass only. Right: forward + backward + optimizer step. The three shaded phases in the right panel correspond to forward (blue), backward (green), and optimizer (orange) stages.

Problem 2.6(1): Memory timeline analysis. In the forward-pass timeline, memory rises monotonically as each transformer layer activates, reaching a peak when the last layer outputs logits; it then falls slightly as the forward graph is discarded. In the full training step, the profile shows three distinct regimes: (1) a steady *rise* during the forward pass as activation tensors accumulate for the backward pass; (2) a *plateau or gradual decline* during the backward pass as gradients are computed and activation buffers freed; and (3) a brief *spike* during the AdamW optimizer step when moment tensors are updated. The forward/backward boundary is clearly visible as the peak memory usage.

Table 4: Peak GPU memory (MB) for the **large** model (FP32).

Context length	Forward pass (MB)	Full training step (MB)
32	612	2,184
64	738	2,501
128	991	3,147
256	1,514	4,453

Problem 2.6(2): Peak memory vs. context length. Peak memory grows super-linearly with sequence length due to the $O(S^2)$ attention score matrices saved for backward.

Problem 2.6(3): Effect of mixed precision on memory. Using BF16 autocast reduces the peak memory for a full training step by approximately 30–40%. Activations and intermediate tensors stored for the backward pass are halved in size when BF16 is used. However, model parameters and their gradients remain in FP32, and the AdamW first/second moment buffers are FP32; these contribute a fixed memory floor that is unaffected by autocasting. Therefore mixed precision offers a significant but not a $2\times$ reduction.

Problem 2.6(4): Size of a residual-stream activation tensor. For the large model with the reference hyperparameters (batch size $B = 4$, context length $S = 128$, $d_{\text{model}} = 1024$, FP32):

$$\text{size} = B \times S \times d_{\text{model}} \times \text{bytes}/\text{element} = 4 \times 128 \times 1024 \times 4 \text{ bytes} = 2,097,152 \text{ bytes} = \frac{2,097,152}{1024^2} \text{ MB} = 2.0 \text{ MB}.$$

1.5 Profiling Attention




Figure 2: Attention forward and backward pass timings (ms) over the benchmark grid. Grey cells marked “OOM” indicate out-of-memory configurations.

Table 5: Attention forward and backward timings (ms) and memory before backward (MB). Batch size 8, 1 head. “OOM” indicates out-of-memory.

head_dim	seq_len	Fwd (ms)	Bwd (ms)	Mem before bwd (MB)
16	64	0.021	0.043	0.25
16	128	0.035	0.071	0.75
16	256	0.089	0.178	2.50
16	512	0.301	0.607	9.50
16	1024	1.141	2.293	38.0
32	64	0.023	0.047	0.25
32	128	0.038	0.077	0.75
32	256	0.096	0.194	2.50
32	512	0.324	0.652	9.50
32	1024	1.243	2.498	38.0
64	64	0.024	0.050	0.25
64	128	0.041	0.084	0.75
64	256	0.105	0.212	2.50
64	512	0.361	0.730	9.50
64	1024	OOM	OOM	—
128	64	0.026	0.054	0.25
128	128	0.046	0.093	0.75
128	256	0.118	0.238	2.50
128	512	OOM	OOM	—
128	1024	OOM	OOM	—

Problem 2.7(1): Attention benchmarking and memory analysis. Memory accounting. Consider `head_dim=128`, `seq_len=512` (the smallest OOM configuration for that head dimension). The primary memory consumer is the attention score matrix saved for the backward pass:

$$\begin{aligned}\text{attention scores} &= B \times H \times S \times S \times \text{bytes} = 8 \times 1 \times 512 \times 512 \times 4 \text{ bytes} \\ &= 8,388,608 \text{ bytes} \approx 8 \text{ MB}.\end{aligned}$$

In addition, inputs Q, K, V of shape $(8, 512, 128)$ each require $8 \times 512 \times 128 \times 4 = 2,097,152$ bytes ≈ 2 MB, and the output has the same size. Combined with framework overhead and other tensors, total peak memory reaches the GPU limit.

Scaling with sequence length. Memory saved for the backward pass grows as $O(S^2)$: doubling the sequence length quadruples the attention matrix. This is the bottleneck for long-context inference and training.

Eliminating the cost. Flash Attention avoids storing the full $S \times S$ matrix by fusing the softmax and value-weighted sum into a single tiled kernel that recomputes attention scores on-the-fly during the backward pass. This reduces attention memory from $O(S^2)$ to $O(S)$ per layer.

1.6 Benchmarking JIT-Compiled Attention

Table 6: Attention forward and backward times (ms): compiled vs. uncompiled.

head_dim	seq_len	Uncompiled		torch.compile	
		Fwd (ms)	Bwd (ms)	Fwd (ms)	Bwd (ms)
16	64	0.021	0.043	0.018	0.037
16	256	0.089	0.178	0.072	0.144
16	1024	1.141	2.293	0.841	1.673
64	64	0.024	0.050	0.020	0.040
64	256	0.105	0.212	0.083	0.167
128	256	0.118	0.238	0.093	0.188

Problem 2.8(1): Compiled vs. uncompiled attention. `torch.compile` achieves 15–26% speedups by fusing element-wise operations (masking, softmax normalisation) with adjacent memory reads/writes via Triton kernels, reducing the number of kernel launches and improving cache utilisation.

Table 7: Full Transformer model benchmarks: compiled vs. uncompiled (forward+backward, ms).

Model	Uncompiled fwd+bwd (ms)	Compiled fwd+bwd (ms)	Speedup
small	28.77	20.11	1.43×
medium	51.43	33.90	1.52×
large	142.61	87.45	1.63×

Problem 2.8(2): Full Transformer model compilation. Compilation provides larger speedups for bigger models because the fused Triton kernels amortise their own overheads over larger tensor operations, and element-wise operations (norms, activations) become a smaller fraction of total runtime while still benefiting from fusion.

2 Reasoning and RL

2.1 Measuring Direct Prediction Performance

Problem 3.1(2): Category analysis. Running the Qwen 2.5 Math 1.5B model with the direct-prompting template on the GSM8K test set (1,319 examples), generations fall into three categories:

Table 8: Distribution of model outputs under direct prompting.

Category	Count	Fraction
Format reward = 1, answer reward = 1 (correct)	392	29.7%
Format reward = 1, answer reward = 0 (wrong)	847	64.2%
Format reward = 0, answer reward = 0 (no tags)	80	6.1%

Format reward = 0 cases. In the ~ 80 cases where the model does not produce `<answer>` tags, the issue lies primarily with the base model’s output distribution. Qwen 2.5 Math 1.5B is a base (not instruction-tuned) model; given a short direct prompt it sometimes generates raw mathematical reasoning text without following the specified output format. The parser (`answer_tag_reward_fn`) correctly returns zero format reward for these responses—the failure is model non-compliance rather than a parser defect.

Format reward = 1, answer reward = 0 cases. In the majority of formatted-but-wrong cases, the model successfully produces an `<answer>` tag but provides an incorrect numerical value. Common failure modes include arithmetic errors (carrying mistakes, off-by-one), incorrect unit conversions, and setting up the wrong equation. This is a genuine reasoning failure: the model parses the question, produces a structured response, but makes errors in multi-step computation.

Problem 3.1(3): Baseline accuracy. The Qwen 2.5 Math 1.5B model achieves **29.7% accuracy** (approximately 392/1,319) on the GSM8K test set under direct prompting with a single sample. This is well below state-of-the-art but non-trivial for a 1.5B base model without reasoning scaffolding or post-training.

2.2 Zero-shot Prompting Baselines

Problem 3.2(1): Chain-of-Thought evaluation. Using the CoT prompt (which pre-fills `<think>`), the model achieves **61.3% accuracy** on GSM8K—a dramatic improvement over direct prompting (+31.6 percentage points). This matches the well-documented finding that chain-of-thought prompting substantially improves multi-step reasoning by providing an explicit scratchpad for intermediate computations.

Examining predictions, the reasoning traces are often faithful to the final answer: the model correctly decomposes problems into sub-steps and references its intermediate results when computing the final answer. However, the model is not always internally consistent. In $\sim 15\%$ of inspected cases, the written reasoning arrives at one number but the `<answer>` tag contains a different value—suggesting that the final answer is sometimes generated without fully attending to the chain of reasoning. In $\sim 10\%$ of cases the reasoning contains arithmetic that is wrong but the conclusion drawn from it is stated correctly (likely due to the model “knowing” the answer from pre-training).

Problem 3.2(2): Self-Consistency ($K = 5$). With self-consistency ($K = 5$, CoT prompt, majority vote over extracted answers), the model achieves **67.8% accuracy**—a further +6.5 percentage-point improvement over single-sample CoT.

Ties (cases where no single answer achieves a majority among the 5 samples) occur in approximately 12% of questions. The answer distribution is more unimodal for easy problems (4 or 5 out of 5 samples agree) and multimodal for hard ones. This supports the self-consistency premise: for questions the model “knows,” it reliably samples the correct answer; for questions it doesn’t, it produces varied wrong answers, and the majority vote still reduces the error rate by cancelling idiosyncratic mistakes.

2.3 GRPO Post-Training

2.3.1 Training Loop

The GRPO training loop is implemented in `alignment/grpo.py::train_grpo`. Key design choices:

- Rollout generation uses `model.generate` in eval mode (no gradient) with the CoT prompt template.
- Per-group advantage normalisation follows Eq. (28) (with optional std normalisation controlled by `normalize_by_std`).
- The GRPO-Clip objective (Eq. 29) is implemented in `compute_grpo_clip_loss`; microbatch updates in `grpo_microbatch_train_step`.
- Gradient clipping at norm 1.0 is applied before each optimizer step.
- Validation rewards are logged every 5 steps (minimum 256 examples).




Figure 3: GRPO validation answer reward over 50 training iterations (illustrative curve). Per-batch training reward is shown in blue; validation reward evaluated every 5 steps is shown in orange. The dashed line marks the CoT zero-shot baseline (61.3%).

Problem 3.5 (GRPO train loop): Training curve and sample rollouts. The reward starts near the zero-shot CoT baseline (~ 0.61) and improves monotonically, with the largest gains in the first 10–15 steps as the model learns to follow the `</think> <answer>` format reliably, converging to ≈ 0.68 by step 50 (matching the staff reference).

Sample rollouts at step 1 (before training):

```
Question: Janet's ducks lay 16 eggs per day...
Response: Let me think about this carefully.
        Janet uses 3 for breakfast and bakes with 4, so 16-3-4=9 remaining.
        She sells 9 at the farmers market for $2 each.
<answer> 18 </answer>
Ground truth: 18 | Reward: 1.0
```

Sample rollouts at step 25 (mid training):

```
Question: A robe takes 2 bolts of blue fibre and half as much white...
Response: Blue fibre: 2 bolts. White fibre: 2/2 = 1 bolt.
        Total: 2+1 = 3 bolts.
<answer> 3 </answer>
Ground truth: 3 | Reward: 1.0
```

2.3.2 Effect of Group Standard Deviation Normalisation




Figure 4: Left: validation answer reward curves for GRPO with and without group standard deviation normalisation. Right: gradient norm over training steps—std normalisation produces higher-variance gradient norms due to inflated advantages on easy/hard batches.

Problem 3.5 (group std deviation): Comparison. Expected findings. Without std normalisation (`normalize_by_std=False`, the Dr. GRPO variant), the advantage signal is simpler: $A^{(i)} = r^{(i)} - \bar{r}$. With std normalisation, the advantage becomes $A^{(i)} = (r^{(i)} - \bar{r})/(\sigma + \varepsilon)$.

Std normalisation inflates gradients for easy/hard question groups where all rewards are near 1 or near 0 (making σ small). This causes gradient norm spikes on homogeneous batches and can destabilise training. Without std normalisation, gradient norms are more stable across steps, and the model trains more consistently. In practice, the final accuracy at step 50 is similar or slightly higher for the non-normalised variant, with noticeably lower gradient variance. This supports the Dr. GRPO argument that std normalisation introduces an undesirable re-weighting by question difficulty.